

HDFS Fault-Tolerance

Thursday, January 26, 2023 3:22 PM

- DFS Failures

• Common Failures

- Entire Node (HW) K-safety
- Partial HW (Non-Storage) K-safety, node-down
- Storage HW RAID, hot-swaps
- Memory/Data Corruption ECC, checksums
- OS Error monitoring, diag/patch, node-down
- Middleware Error hdfs/hadoop, log/debug

• HDFS Failures

- Data Node

- Heartbeat
- Mark as "dead" on lost heartbeat
 - ↳ re-replicate affected blocks

- Name Node

- Persist state of Meta Data
 - fsimage
 - edit log} RDBMS

• Manual Failover (HA)

- Primary (Hot) X
- Secondary (Cold) ✓ rm failover

* Hot-Hot / Hot-Warm (Multi-Master)

* Microservices: state replicated across (min 2)
Name Nodes via HTTP/REST

* Zookeeper: Maintains Primary/Secondary failover

* Usage: CLI

hdfs dfs ← old
hadoop fs ← new (+object store)
Azure, S3...

