Map-Reduce Framework

Tuesday, February 14, 2023 2:21 PM

- Map-Reduce Framework
  - Framework
    - Services: (Zookeeper) HDFS / YARN
    - Interfaces: (Common), API

Wrapper Types
Driver Hooks / Callbacks

* Data Exchange: Framework ↔ User functions
  - Key-Value Pairs (KVPs)

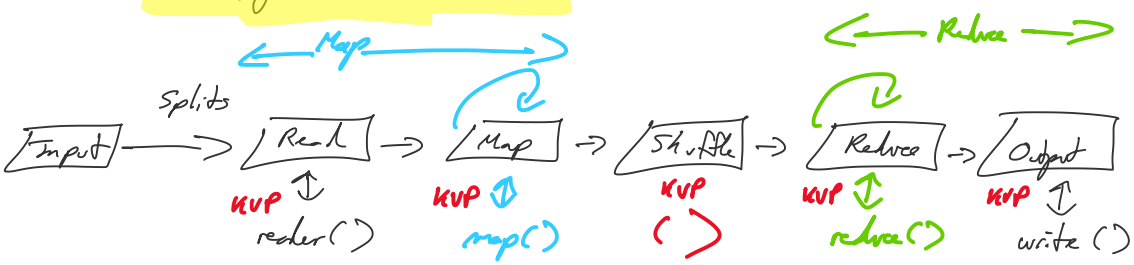

Ex.

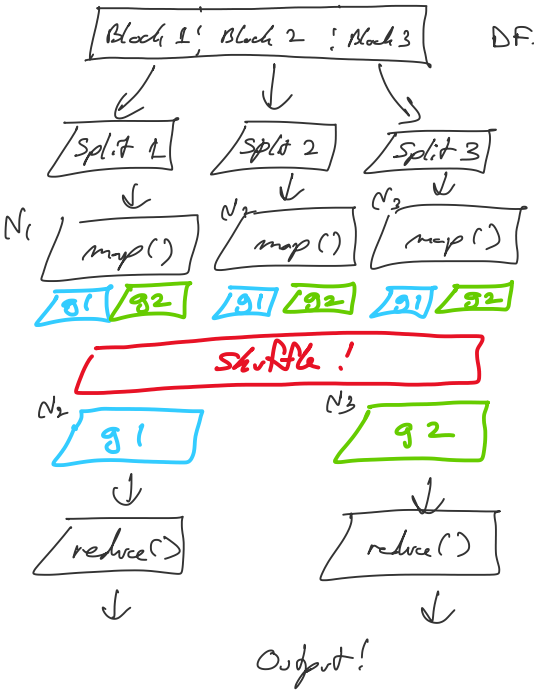

* For this example, Blocks & Splits Align!

Batch Example:

60616, 3
60616, 2
60617, 5
60617, 2

* KVPs: Can be any POD (Plain-Old-Datatype) OR any Complex Datatype
  - Struct
  - Array/List
  - BitMarkers ...
  - ⋮

Ex. Name = "Jay"

User = {

Name = "Ja" ? can also

id = 401 } ← be a key!
}

* Languages: Java → C/C++...
Scala
→ Python, R, ... (SQL)

- Customizations
  - Map Hooks

Input → Map → Combine → Partition:
Record Reader — Mapper: - Init - Final — Combiner — Partitioner

- Reduce Hooks

Shuffle → Reduce → Output
Combiner — Reducer: -Init -Final — Output Formatter

* Combiners
- Localized Reducer → to minimize shuffle I/O cost
- Not guaranteed to execute!
- Mapper has a memory buffer
  - on overflow → spills to disk: Spill Files
  - Run Combiner to merge spill files!

* Partitioner
- Intermediate KVP → split/Sharded to have one shard per reducer
- Hash object (md5, sha1..)
  - hash(key) % # of reducers
- Keyspace of intermediate KVP → Distribute evenly across Reducer nodes, need to guarantee mapping to identical Reducer nodes

- Shuffle & Sort
  - Map
    - Output to buffer, spill to disk
    - Run Combiner if needed

- Single partitioned file for Reduce step
  ↑ sorted keys within each partition

- Reduce
  - Copy map files by partition to Reduce nodes
  - Run Combiner if needed
  - Multi-Pass Merge
    Sort all Intermediate KVP for Reducer
  - Reducer is run